

## Homework 4. Supervised learning

In Lab 4 you will use supervised learning to solve a real-world classification problem. To prepare for the lab, in the homework you will go over a toy problem that you can run by hand.

*Logistic regression* consists of estimating the probability of each of two actions,  $\mathcal{A} = \{0, 1\}$ , given a set of examples,  $\mathcal{D} = \{(x_1, a_1), \dots, (x_N, a_N)\}$ , with  $a_n \in \mathcal{A}, n = 1, \dots, N$  and where each state  $x_n$  is described by a number of features  $\phi_1, \dots, \phi_K$ . In logistic regression, we assume that

$$\pi(1 \mid x) \stackrel{\text{def}}{=} \mathbb{P}[a = 1 \mid x = x] = \frac{1}{1 + e^{-(\mathbf{w}^\top \boldsymbol{\phi}(x) + w_0)}}, \quad (1)$$

where  $\mathbf{w}$  and  $w_0$  are the parameters to be learned.

Training logistic regression consists of finding the parameters  $\mathbf{w}, w_0$  that minimize the *negative log likelihood* of the data, i.e.,

$$J(\pi) = -\log \prod_{n=1}^N \pi(a_n \mid x_n), \quad (2)$$

which can be done, for example, using gradient descent.

Consider the following dataset, comprising 5 points described by two attributes,  $\phi_1$  and  $\phi_2$ .

$\phi_1(x)$	$\phi_2(x)$	$a$
1.0	3.0	0
2.0	2.0	1
2.0	1.0	0
3.0	1.0	0
5.0	4.0	1

### Exercise 1.

- (a) Plot the training data provided. Is the data linearly separable? Why?
- (b) Initializing the weights of the logistic regression classifier to 0, run 1 iteration of gradient descent with a step-size  $\alpha = 1$ , and indicate the resulting weights.
- (c) Indicate the equation defining the decision boundary corresponding to the weights obtained after the update in Question (b), and indicate it in the plot of Question (a). Does it properly classify the points in the training set?